

# Numerical Linear Algebra: Class Project Mar 22

John Kallaugher

March 21, 2016

## 1 IMPLEMENTATION

I chose  $n_r = m_r = 4$  based on the length of the vector values we were using. I then implemented a rank-1 update for length 4 vectors by using a butterfly permute executed with Intel intrinsic functions. (using ICC as my compiler) This update was executed in a loop that executes  $k_c$  times, each time adding the results of the rank-1 update to the matrix calculated so far.

I found this gave approximate a 4-fold increase in performance over the naïve C implementation, and about  $\frac{1}{6}$  of the performance of the reference implementation ( $\frac{1}{8}$  of peak.) Experimenting with changing the values of the DGEMM constants (other than  $n_r$  and  $m_r$ ) did not seem to have any effect on performance.

## 2 PERFORMANCE GRAPH

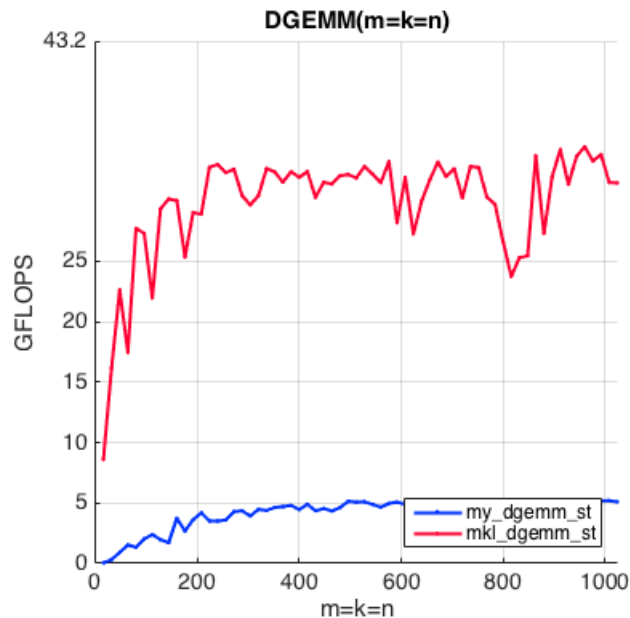


Figure 2.1: Single-threaded performance of my\_dgemm against the reference implementation.